

Comparison of Human Chest Diameter and Height

Code ▾

Hide

```
library(dplyr)
```

Registered S3 method overwritten by 'dplyr':

```
method      from  
print.rowwise_df
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

```
filter, lag
```


The following objects are masked from 'package:base':

```
intersect, setdiff, setequal, union
```


Hide

```
bdims_filtered <- select(bdims_csv, che.di, hgt)
```

```
plot(che.di ~ hgt, data = bdims_filtered)
```


[Hide](#)

```
bdims_filtered$che.di %>% hist(main="Histogram of Chest Diameter")
```




[Hide](#)


```
bdims_filtered$hgt %>% hist(main="Histogram of Height")
```

[Hide](#)

```
plot(log(che.di) ~ log(hgt), data = bdims_filtered)
```

[Hide](#)

```
plot(sqrt(chest.di) ~ sqrt(hgt), data = bdims_filtered)
```



Start linear regression model fitting test H_0 Data does not fit H_a Data fits

Hide

```
lmtest <- lm(log(che.di) ~ log(hgt), data = bdims_filtered)
lmtest %>% summary()
```

Call:

```
lm(formula = log(che.di) ~ log(hgt), data = bdims_filtered)
```

Residuals:

Min	1Q	Median	3Q	Max
-0.22123	-0.05121	0.00044	0.05057	0.22962

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	-2.42566	0.31592	-7.678	8.43e-14 ***
log(hgt)	1.11888	0.06145	18.208	< 2e-16 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.0759 on 505 degrees of freedom

Multiple R-squared: 0.3963, Adjusted R-squared: 0.3951

F-statistic: 331.5 on 1 and 505 DF, p-value: < 2.2e-16

Hide

```
lmtest %>% confint()
```

	2.5 %	97.5 %
(Intercept)	-3.0463489	-1.804974
log(hgt)	0.9981551	1.239607

P value is very small as $p < 0.001$. The model is statistically significant F Statistic is $F(1,505) = 331.5$

Model coefficients are as follows

alpha (intercept) = -2.42566 == -2.43 beta (slope) = 1.11888 == 1.12

The slope represents the average change in y following a one unit increase in x. As log(hgt) increases by 1 unit, log(che.di) changes on average by 1.12

Not required to perform as the p value is < 0.001 (Hypothesis testing for the model coefficients)

Intercept H_0 alpha = 0 H_a alpha is not equal to 0

Slope H_0 beta = 0 H_a beta is not equal to 0)

Error residual is as follows

$y = y_i - \hat{y}_{cap}$

y_i for che.di = 30.3 is $\log(30.3) = 3.4111$

\hat{y}_{cap} (predicted value) for che.di = 30.3 and hgt = 180.0 3.3861

hence $y = 3.4111 - 3.3861 = 0.025$

Testing assumptions :-


Independence - Proved because of the data

Linearity - Not perfectly linear (plot 1) Normality of residuals - Proved from plot 2 (most points lie on the black line)


Homoscedasticity - Proved from plot 3 as the red line is almost flat

[Hide](#)


```
lmtest %>% plot(which = 1)
```

[Hide](#)

```
lmtest %>% plot(which = 2)
```


[Hide](#)

```
lmtest %>% plot(which = 3)
```




Hide

```
lmtest %>% plot(which = 4)
```



Hide

```
lmtest %>% plot(which = 5)
```





Residual vs Leverage

In the example below, no values fall close to 0.5 band. In fact, the bands are not even visible

[Hide](#)

```
lmtest %>% plot(which = 5)
```



Linear Regression - Strength and Direction of Linear Relationships

Correlation coefficient

Hide

```
r <- cor(log(bdims_filtered$che.di), log(bdims_filtered$hgt), use = "complete.obs")
```

```
r
```

```
[1] 0.6295466
```

CI for Correlation coefficient

Hide

```
install.packages("psychometric")
```

WARNING: Rtools is required to build R packages but is not currently installed. Please download and install the appropriate version of Rtools before proceeding:

<https://cran.rstudio.com/bin/windows/Rtools/>

Installing package into 'C:/Users/smart/OneDrive/Documents/R/win-library/3.6'
(as 'lib' is unspecified)
also installing the dependency 'multilevel'

trying URL 'https://cran.rstudio.com/bin/windows/contrib/3.6/multilevel_2.6.zip'
Content type 'application/zip' length 433488 bytes (423 KB)
downloaded 423 KB

trying URL 'https://cran.rstudio.com/bin/windows/contrib/3.6/psychometric_2.2.zip'
Content type 'application/zip' length 194016 bytes (189 KB)
downloaded 189 KB

package 'multilevel' successfully unpacked and MD5 sums checked
package 'psychometric' successfully unpacked and MD5 sums checked

The downloaded binary packages are in
C:\Users\smart\AppData\Local\Temp\RtmpgT0gno\downloaded_packages

[Hide](#)

```
library(psychometric)
```

Loading required package: multilevel
Loading required package: nlme

Attaching package: 'nlme'

The following object is masked from 'package:dplyr':

collapse

Loading required package: MASS

Attaching package: 'MASS'

The following object is masked from 'package:dplyr':

select

[Hide](#)

```
CIr(r = r, n = length(bdims_filtered$hgt), level = .95)
```

```
[1] 0.5739283 0.6793837
```

[Hide](#)

```
detach("package:psychometric", unload=TRUE)
```

Simple linear regression summary:

Linearity was assumed, normality of residuals OK, homoscedasticity OK, no influential cases. $r = 0.6296$, $r_sqr = 0.3964$ Model ANOVA, $F(1,505) = 331.5$, $p < 0.001$ $a = -2.43$, $p < 0.001$, 95% CI (-3.0463489 -1.804974) $b = 1.12$, $p < 0.001$, 95% CI (0.9981551 1.239607)

Decision:

Overall model: Reject H_0 . Intercept: Reject H_0 . Slope: Reject H_0 . $\log(\text{chest diameter}) = -2.43 + 1.12\log(\text{Height})$

What do we conclude?: There was a statistically significant positive linear relationship between chest diameter and height.

Making predictions Take a random example from the dataset and make predictions, refer to the slide

[Hide](#)

```
bdims_filtered %>% summarise(Min = min(hgt,na.rm = TRUE),
  Q1 = quantile(hgt,probs = .25,na.rm = TRUE),
  Median = median(hgt, na.rm = TRUE),
  Q3 = quantile(hgt,probs = .75,na.rm = TRUE),
  Max = max(hgt,na.rm = TRUE),
  Mean = mean(hgt, na.rm = TRUE),
  SD = sd(hgt, na.rm = TRUE),
  n = n(),
  Missing = sum(is.na(hgt)))
```

Min <dbl>	Q1 <dbl>	Median <dbl>	Q3 <dbl>	Max <dbl>	Mean <dbl>	SD <dbl>	n <int>	Missing <int>
147.2	163.8	170.3	177.8	198.1	171.1438	9.407205	507	0

1 row

[Hide](#)

```
bdims_filtered %>% summarise(Min = min(che.di,na.rm = TRUE),
  Q1 = quantile(che.di,probs = .25,na.rm = TRUE),
  Median = median(che.di, na.rm = TRUE),
  Q3 = quantile(che.di,probs = .75,na.rm = TRUE),
  Max = max(che.di,na.rm = TRUE),
  Mean = mean(che.di, na.rm = TRUE),
  SD = sd(che.di, na.rm = TRUE),
  n = n(),
  Missing = sum(is.na(che.di)))
```

Min <dbl>	Q1 <dbl>	Median <dbl>	Q3 <dbl>	Max <dbl>	Mean <dbl>	SD <dbl>	n <int>	Missing <int>
22.2	25.65	27.8	29.95	35.6	27.97377	2.74165	507	0

1 row